

# Lab 2 Write-Up

---

Jordi Burbano (204 076 325), Keisuke Daimon (604 547 017)

February 17, 2015

## I. DESIGN DECISIONS

The `BufferPool` class uses an LRU eviction policy. This is implemented using a `LinkedList<Integer>` of the hash codes that are the keys to the `cachedPages` hash map (which maps int hash codes to `Page` objects); these hash codes are listed from the least recently used page at the head to the most recently used page at the end. Updating the linked list occurs when a page is referenced in `BufferPool.getPage()` so that it moves to the end of the linked list as the most recently page. It also occurs during `BufferPool.evictPage()` when the evicted page's hash code is removed from the linked list. `BufferPool.evictPage()` simply refers to the hash code at the head of the linked list to remove the corresponding page from the buffer pool's hash map.

## II. CHANGES TO THE API

We made no changes to the API.

## III. MISSING OR INCOMPLETE ELEMENTS OF OUR CODE

For the methods `getJoinField1Name()` and `getJoinField2Name()` in `Join.java`, we are able to access a table's name by consulting `Catalog.getTableName()`. However, we did not know how to access the alias name as described by the requirement that field names should be quantified by alias or table name.

## IV. LOGISTICS

We spent approximately 25 man-hours on the project.

We spent a significant amount of time trying to pass the `EvictionTest` system test. First, we discovered an issue attributable to not closing `RandomAccessFile` objects when reading or writing pages. Afterwards, much time was spent on overcoming an error that stated that 80 MB of RAM were being used when the limit was 5 MB. This error was related to our previous implementation of `TransactionFileDbIterator`, which stored all the pages in an `ArrayList<Page>`. To consume less memory, `TransactionFileDbIterator` now stores an `ArrayList<PageId>`; this change reduced the consumption of RAM to 1 MB.